

# Machine Learning

## A Brief Introduction

Marvin N. Wright

Leibniz Institute for Prevention Research & Epidemiology – BIPS  
University of Bremen  
University of Copenhagen

December 2022

# Outline

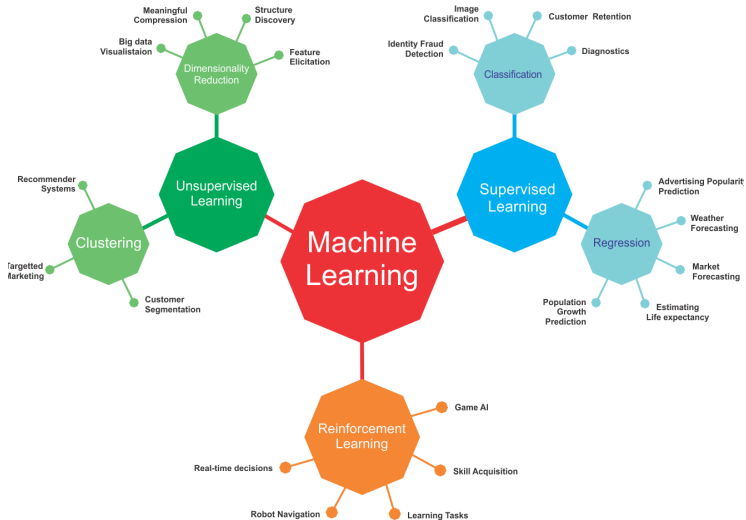
---

2

1. Introduction
2. Supervised Learning
3. Decision Trees & Random Forests
4. Model Evaluation & Resampling
5. Penalized Regression
6. Artificial Neural Networks
7. Hyperparameter Tuning & Benchmarking
8. Discussion

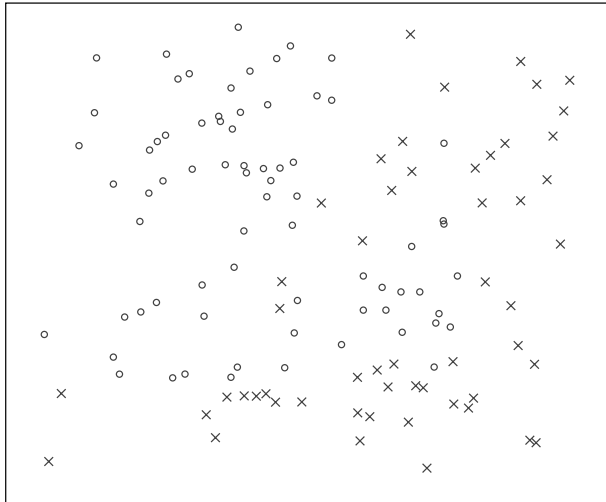
# Machine Learning

3

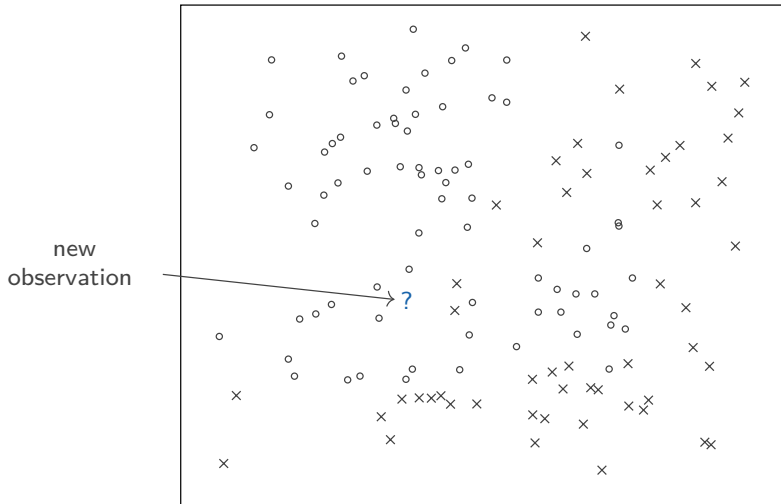


# k-Nearest Neighbors


---



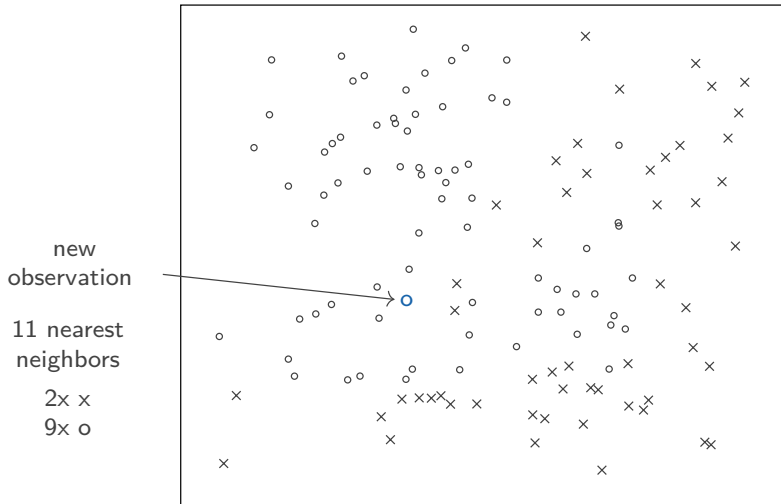
# k-Nearest Neighbors



# k-Nearest Neighbors



# k-Nearest Neighbors



# Outline

---

5

1. Introduction
2. Supervised Learning
3. Decision Trees & Random Forests
4. Model Evaluation & Resampling
5. Penalized Regression
6. Artificial Neural Networks
7. Hyperparameter Tuning & Benchmarking
8. Discussion



## Example: House Prices

Predict the price for a house in a certain area

6

Features $x$				Target $y$
square footage of the house	number of bedrooms	swimming pool (yes/no)	...	house price in US\$
1,180	3	0	...	221,900
2,570	3	1	...	538,000
770	2	0	...	180,000
1,960	4	1	...	604,000



# Example: Length of Hospital Stay

Predict days a patient has to stay in hospital

7

Features $x$					Target $y$
diagnosis category	admission type	gender	age	...	Length-of-stay in the hospital in days
heart disease	elective	male	75	...	4.6
injury	emergency	male	22	...	2.6
psychosis	newborn	female	0	...	8
pneumonia	urgent	female	67	...	5.5



# Example: Life Insurance

Predict risk category for a life insurance customer

8

Features $x$				Target $y$
job type	age	smoker	...	risk group
carpenter	34	1	...	3
stuntman	25	0	...	5
student	23	0	...	1
white-collar worker	39	0	...	2



# Supervised Learning

9

Learn a functional relationship between **features**  $x$  and **target**  $y$

Features $x$		Target $y$
People in Office (Feature 1) $x_1$	Salary (Feature 2) $x_2$	Worked Minutes Week (Target Variable)
4	4300 €	2220
12	2700 €	1800
5	3100 €	1920

$n = 3$

$p = 2$

$x_1^{(2)}$

$x_2^{(1)}$

$y^{(3)}$

# Supervised Learning

10

Use labeled data to learn a model  $f$

Use model  $f$  to predict target  $y$  of new data



# Supervised Learning


## Model

Functional relationship between **features**  $x$  and **target**  $y$

## Learner (or inducer)

Algorithm for finding model

11



# Supervised Learning

---

12

## Example

- Learner: Artificial neural network (as a concept)
- Model: Actual network with learned weights

## Models differ in size and complexity

- Linear model: Coefficients  $\beta$
- Neural network: Weights for all units in all layers
- Decision trees: Many binary splits
- $k$ -nearest neighbors: Complete training data

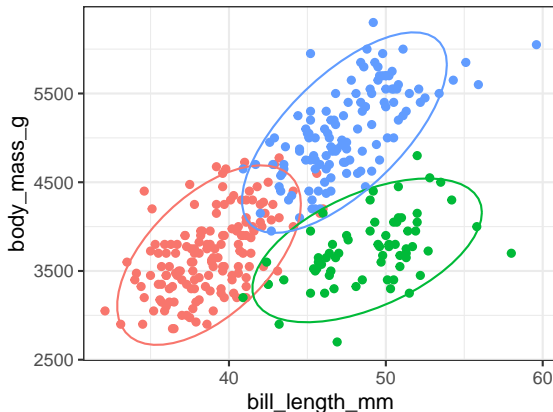
# Supervised Learning

## Unsupervised Learning

No **target**  $y$  available

Search for patterns in the data  $x$ , e.g. clustering:

13





# Outline

---

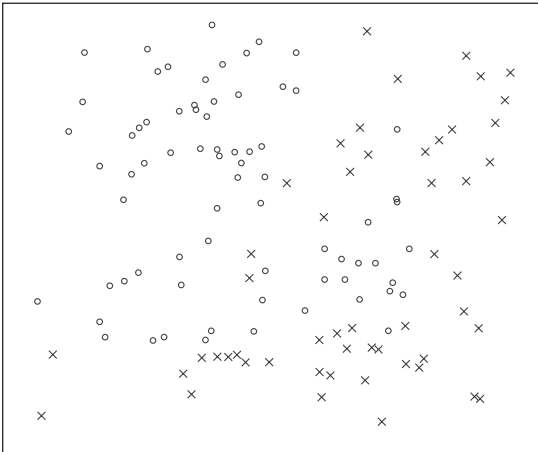
14

1. Introduction
2. Supervised Learning
3. Decision Trees & Random Forests
4. Model Evaluation & Resampling
5. Penalized Regression
6. Artificial Neural Networks
7. Hyperparameter Tuning & Benchmarking
8. Discussion

# Decision Trees

---

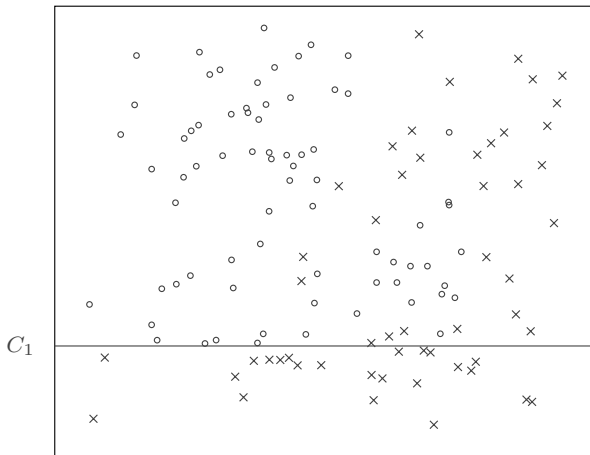
15



# Decision Trees


---

15

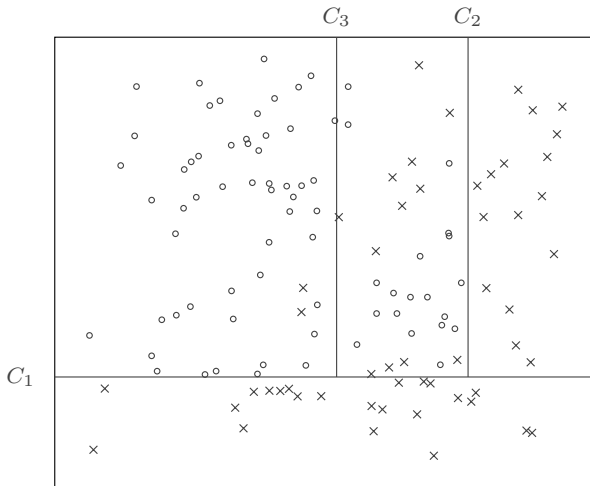


# Decision Trees

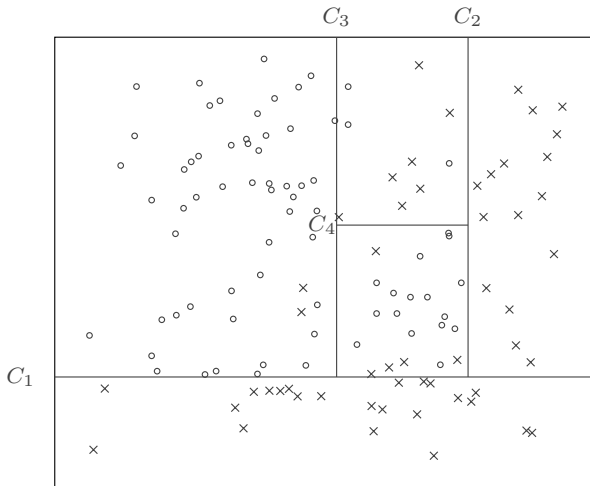
---




# Decision Trees



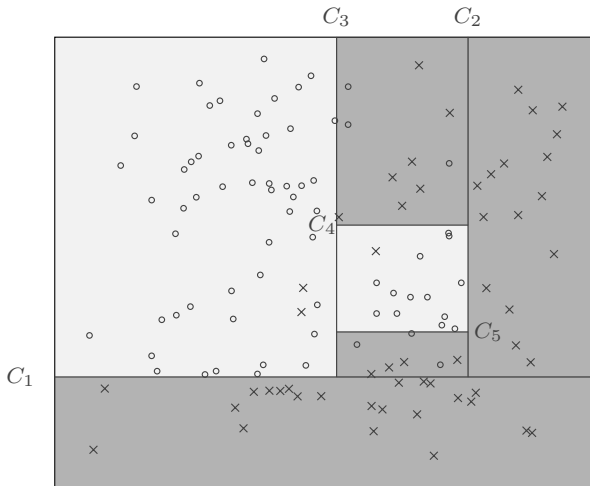
# Decision Trees



# Decision Trees





# Decision Trees


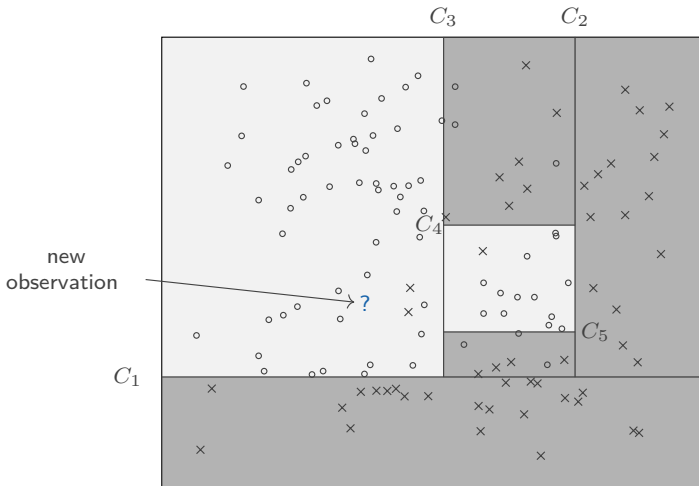




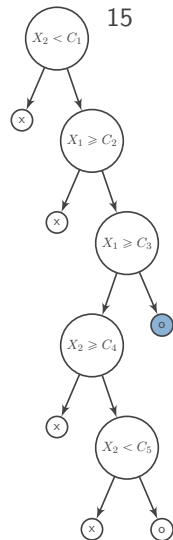
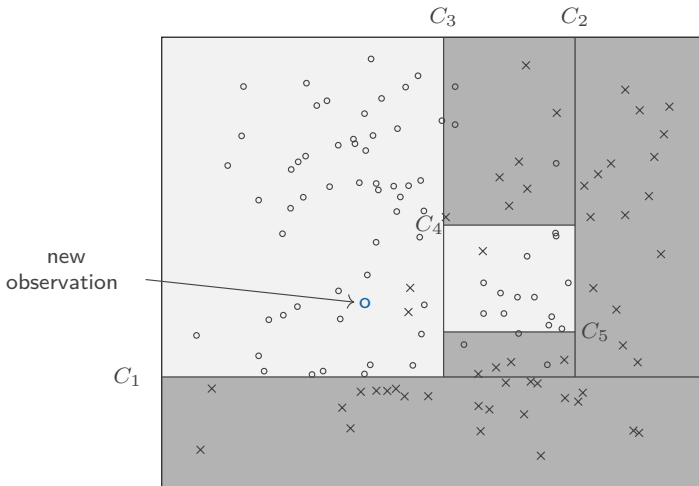
# Decision Trees



# Decision Trees



# Decision Trees



# Decision Trees

---

16

## Advantages of decision trees

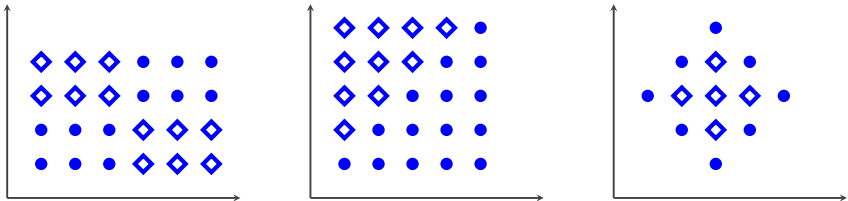
- Procedure intuitive
- Small trees simple to interpret
- Intrinsic variable selection
- Simple handling of outliers
- Fast training
- Usually better prediction performance than kNN

# Decision Trees

17

## Disadvantages of decision trees

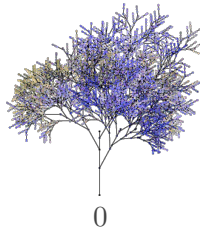
- Trees unstable
- Pruning can be computationally intensive
- Usually worse prediction performance than random forests (covered later) and boosted trees
- Problematic data sets



# Random Forests

---


18



# Random Forests

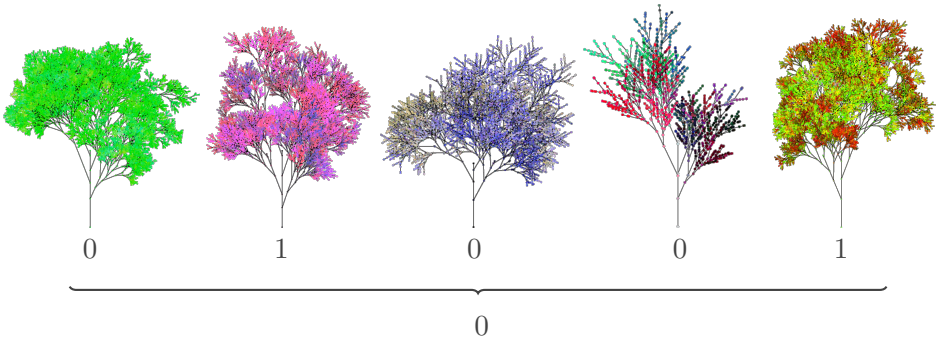
---

18



# Random Forests

18

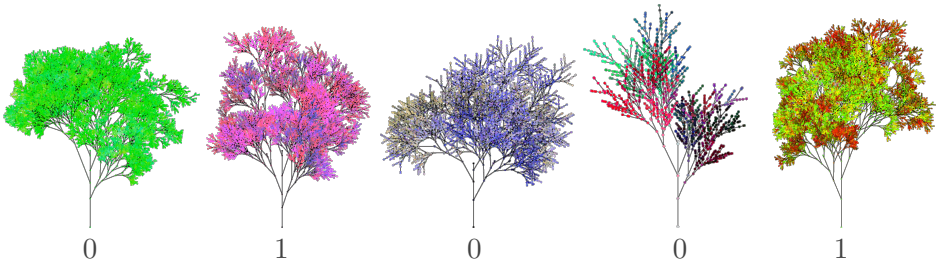


Classification: **majority vote** over all trees



# Random Forests

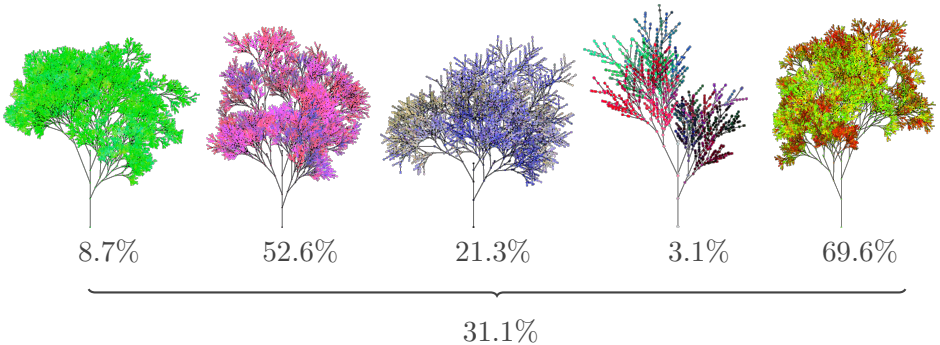
18



Classification: **majority vote** over all trees  
Identical to average over all trees, cut point 0.5

# Random Forests

18



Probability estimation: Average over all trees

# Random Forests

---

19

## Two components of randomization

- Data manipulation in rows: bootstrapping / subsampling
- Data manipulation in columns: feature subsampling

# Random Forests

---

20

## Bootstrap aggregating (bagging)

- Ensemble = committee of experts
- Single weak learner = single committee member
- Ensemble decision = committee decision

Fundamental idea of bagging (bootstrap aggregating)

Any machine can be used as *base learner*, e.g. kNN or tree

# Random Forests

---

21

## Bootstrapping

- Sampling **with** replacement
- Original sample size  $n$ , resampled sample size  $n$
- On average  $\lim_{n \rightarrow \infty} \left(1 - \frac{1}{n}\right)^n \approx 0.632 \approx 2/3$  resampled

## Subsampling

- Sampling **without** replacement
- Original sample size  $n$ , resampled sample size  $< n$
- Standard: resampling of  $0.632n$

# Random Forests

---

22

## Feature subsampling

At a node consider only subset of features

- Trees vary
- “Experts” differ in their opinion
- Reduce correlation between trees

## Number of features considered at split

$m_{\text{try}} = \sqrt{d}, \ln d \text{ or } d/3 \rightarrow \text{Tuning possible (later)}$

# Random Forests

---

23

## Random forest algorithm

For each tree

1. Draw bootstrap sample with replacement
2. Grow tree
  - a) Use random subset of variables ( $m_{try}$ ) at each node
  - b) Stop if minimum node size reached
3. Determine proportion of '1' in each terminal node

New subject

1. Drop down subject in each single tree
2. Store proportion from all trees
3. Average proportion of '1's over all trees

# Random Forests

---

24

## Advantages of random forests

- As with trees: Procedure intuitive, intrinsic variable selection, simple handling of outliers, fast training
- Work well with high dimensional data
- Work well without (or with only a little) tuning
- Usually better prediction performance than a single tree



# Random Forests

---

25

## Disadvantages of random forests

- Not simple to interpret
- Sometimes worse prediction performance than well tuned boosted trees
- Bad prediction performance on image, text and speech data

# Outline

---

26

1. Introduction
2. Supervised Learning
3. Decision Trees & Random Forests
4. Model Evaluation & Resampling
5. Penalized Regression
6. Artificial Neural Networks
7. Hyperparameter Tuning & Benchmarking
8. Discussion

# Model Evaluation

---

27

## How good is a prediction model?

Compare true target  $y$  with predicted target  $\hat{y}$

### Examples

- How many patients correctly diagnosed?
- How many emails correctly detected as ham or spam?
- How close is the predicted price of a house to the true value?
- How close is the length of hospitalization to the true value?

# Model Evaluation

---

28

## Dichotomous (binary) outcome

- Proportion of correct classifications (PC); also accuracy:

$$\widehat{PC} = \frac{1}{n} \sum_{i=1}^n \mathbf{1}_{y_i = \hat{y}_i}$$

- Sensitivity, specificity, ROC, AUC:  $\hat{\mathbb{P}}(y = 1 \mid x)$
- Brier score (BS), i.e., MSE of probability estimates; also probability score (PS):  $\widehat{BS} = \frac{1}{n} \sum_{i=1}^n \left( y_i - \hat{\mathbb{P}}(y_i = 1 \mid x_i) \right)^2$

## Multicategory outcome

- Proportion of correct classifications (PC)
- Averaged class-wise PC
- ROC, AUC: several extensions

# Model Evaluation

---

29

## Continuous outcome

- MSE:  $\widehat{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$
- MAE:  $\widehat{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$
- RMSE:  $\widehat{RMSE} = \sqrt{\widehat{MSE}}$
- Explained variance:  $\hat{R}^2 = \frac{1 - \widehat{MSE}}{\widehat{\text{Var}}(y)}$

## Survival outcome

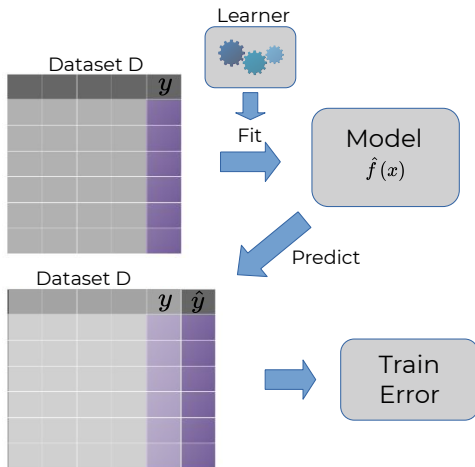
- Time-dependent Brier Score
- Integrated Brier score
- C-Index

# Model Evaluation

## Training error

30

Evaluate performance on training data




# Model Evaluation

## Training error

30

Evaluate performance on training data




**Problem:**  
**Overfitting**

# Model Evaluation

## Overfitting

31

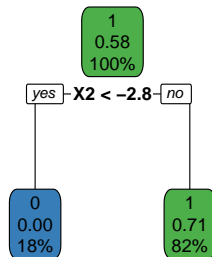
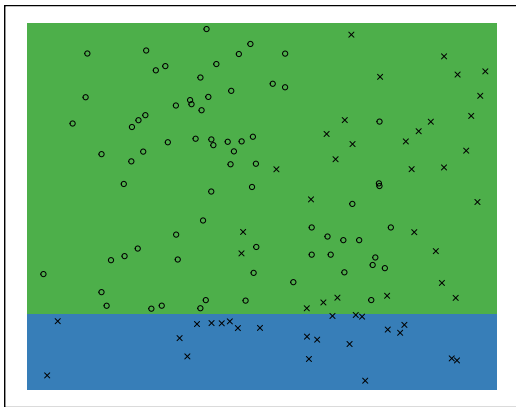




# Model Evaluation

## Overfitting



31



# Model Evaluation

## Overfitting

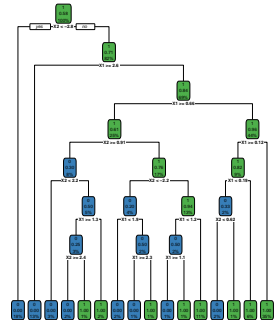
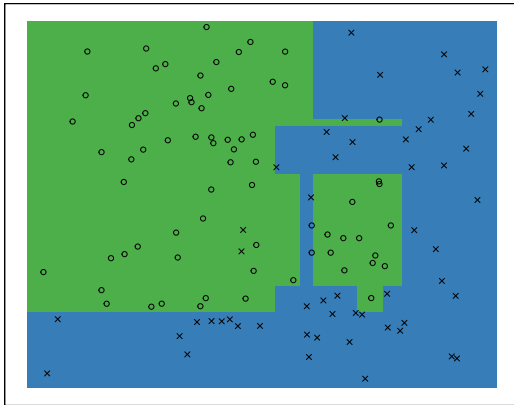
31



# Model Evaluation

## Overfitting

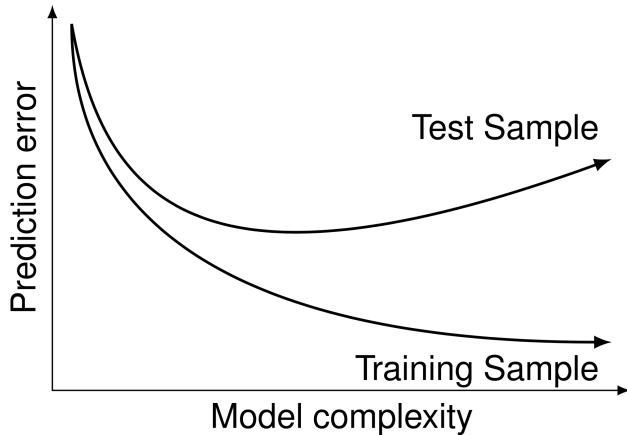
31



# Model Evaluation

32

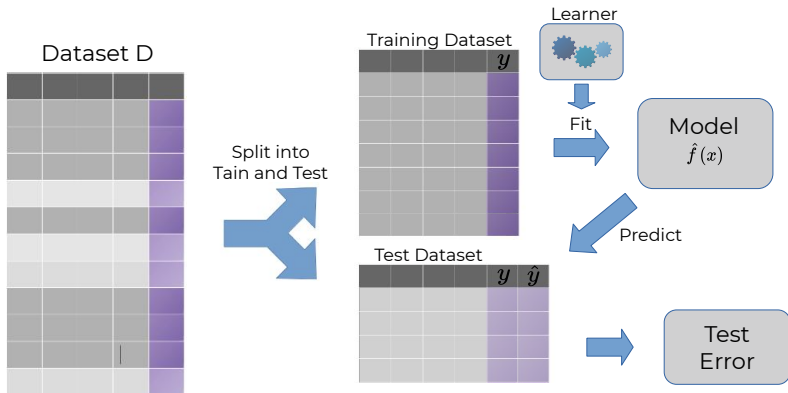
Overfitting



# Model Evaluation

33

## Test error



# Model Evaluation

---

34

## Training and test error

- Training error heavily biased
- Test error (almost) unbiased but variance unknown

## Resampling

- Repeated training/test splits (subsampling)
- Cross validation
- Repeated cross validation
- Bootstrap

# Model Evaluation

---

## Hyperparameters

35

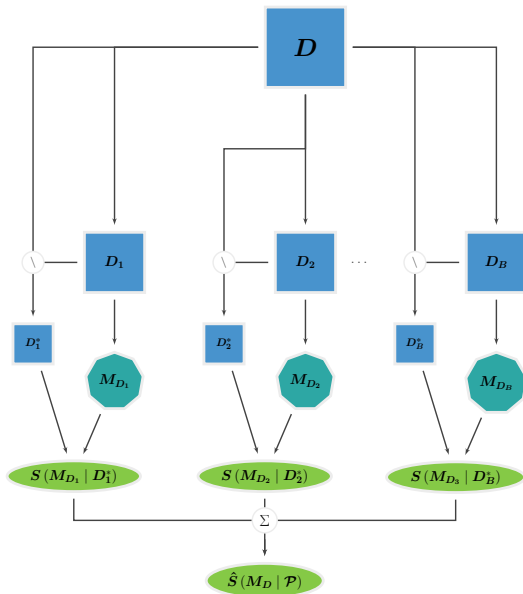
Most (all?) learners have hyperparameters, e.g.:

- $k$ -nearest neighbors: Number of neighbors  $k$ , distance weighting, etc.
- Decision trees: Tree depth, splitting criterion, etc.
- Neural networks: Number and size of layers, activation function, regularization, etc.

## Hyperparameter tuning

- Optimize (tune) the hyperparameters
  - Do not tune and evaluate on same data
- 3-fold split into training, validation, test
- Nested resampling

# Resampling





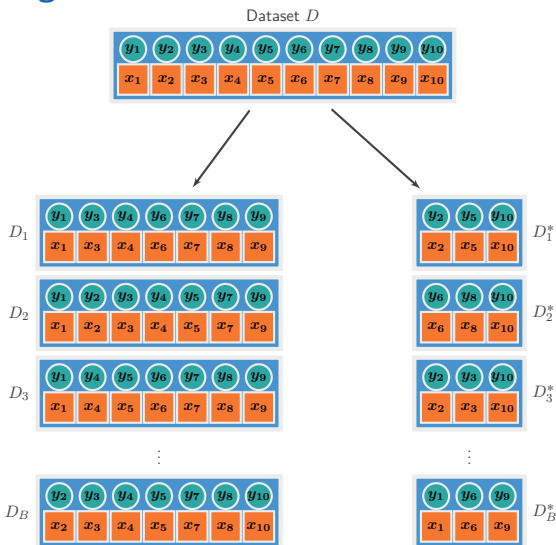
## Resampling

- Estimate performance on independent data
- Used for
  - Performance estimation
  - Hyperparameter tuning
  - Model selection
- Resampling based performance estimation
  1. Split dataset in several (smaller) datasets  $D_b$
  2. On each dataset  $D_b$ :
    - 2.1 Train learner
    - 2.2 Estimate performance on  $D_b^* = D \setminus D_b$
  3. Aggregate performance estimates

# Resampling

## Subsampling

38



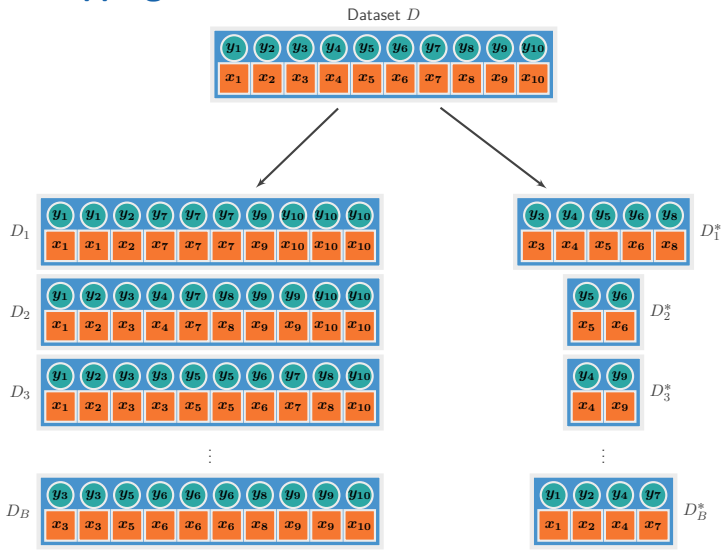
## Subsampling

- Sample  $B$  training datasets  $D_b$  from  $D$  without replacement, usually  $n_b = \frac{2}{3}n$
- Use  $D_b^* = D \setminus D_b$  as test datasets
- $D_b$  and  $D_b^*$  disjunct
- $D_1$  and  $D_2$  not disjunct
- $D_1^*$  and  $D_2^*$  not disjunct
- Performance estimator biased
- No optimal  $B$ , usually  $100 < B < 1000$
- Special case with  $B = 1$ : Single train/test split (holdout)

# Resampling

## Bootstrapping

40



# Resampling

---

41

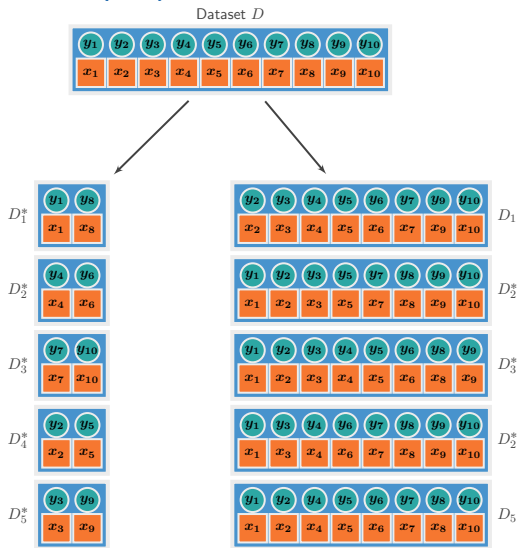
## Bootstrapping

- Sample  $B$  training datasets  $D_b$  from  $D$  with replacement, usually  $n_b = n$
- Use  $D_b^* = D \setminus D_b$  as test datasets
- $D_b$  and  $D_b^*$  disjunct
- $D_1$  and  $D_2$  not disjunct
- $D_1^*$  and  $D_2^*$  not disjunct
- Performance estimator biased
- Adaptive weighting to reduce bias (.632+ bootstrap)
- Small variance (large  $B$ )
- No optimal  $B$ , usually  $100 < B < 1000$

# Resampling

## Cross validation (CV)

42




## Cross validation (CV)

- Split  $D$  in  $B$  test datasets  $D_b^*$
- Use  $D_b = D \setminus D_b^*$  as training datasets
- $D_b$  and  $D_b^*$  disjunct
- $D_1$  and  $D_2$  not disjunct
- $D_1^*$  and  $D_2^*$  disjunct
- Special case with  $B = n$ : Leave-one-out CV (LOOCV)
  - Small bias, high variance
  - Long runtime
- No optimal  $B$ , usually  $B = 5, 10$ 
  - Slightly more bias than LOOCV, but lower variance
  - Lowest  $B$  of all resampling methods  $\rightarrow$  fast computation

# Resampling

44

## How to choose the resampling method?





# Outline

---

45

1. Introduction
2. Supervised Learning
3. Decision Trees & Random Forests
4. Model Evaluation & Resampling
5. Penalized Regression
6. Artificial Neural Networks
7. Hyperparameter Tuning & Benchmarking
8. Discussion

# Penalized Regression

---

46

## Generalized linear model

$$\begin{aligned} g(\mathbb{E}(Y)) &= \beta_0 + \beta_1 \cdot X_1 + \dots + \beta_p \cdot X_p \\ &= X\beta \end{aligned}$$

$g$ : Link function

# Penalized Regression

---

46

## Generalized linear model

$$\begin{aligned} g(\mathbb{E}(Y)) &= \beta_0 + \beta_1 \cdot X_1 + \dots + \beta_p \cdot X_p \\ &= X\beta \end{aligned}$$

$g$ : Link function

## Linear model

$$\mathbb{E}(Y) = X\beta$$

# Penalized Regression

---

47

## Ordinary least squares

Minimize squared differences

$$\begin{aligned} L_{\text{OLS}} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\ &= \|y - X\beta\|_2^2 \\ &= (y - X\beta)'(y - X\beta) \end{aligned}$$

# Penalized Regression

---

47

## Ordinary least squares

Minimize squared differences

$$\begin{aligned}L_{\text{OLS}} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\&= \|y - X\beta\|_2^2 \\&= (y - X\beta)'(y - X\beta)\end{aligned}$$

Solution:

$$\beta_{\text{OLS}} = (X'X)^{-1} X'y$$

# Penalized Regression

---

48

## Ridge regression

Penalize large parameter estimates (L2 regularization)

$$\begin{aligned} L_{\text{Ridge}} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \lambda \sum_{j=1}^m \beta_j^2 \\ &= \|y - X\beta\|_2^2 + \lambda \|\beta\|_2^2 \end{aligned}$$

# Penalized Regression

---

48

## Ridge regression

Penalize large parameter estimates (L2 regularization)

$$\begin{aligned} L_{\text{Ridge}} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \lambda \sum_{j=1}^m \beta_j^2 \\ &= \|y - X\beta\|_2^2 + \lambda \|\beta\|_2^2 \end{aligned}$$

Solution:

$$\beta_{\text{Ridge}} = (X'X + \lambda I)^{-1} X'y$$

# Penalized Regression

---

48

## Ridge regression

Penalize large parameter estimates (L2 regularization)

$$\begin{aligned} L_{\text{Ridge}} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \lambda \sum_{j=1}^m \beta_j^2 \\ &= \|y - X\beta\|_2^2 + \lambda \|\beta\|_2^2 \end{aligned}$$

Solution:

$$\beta_{\text{Ridge}} = (X'X + \lambda I)^{-1} X'y$$

**Shrink parameter estimates towards zero**



# Penalized Regression

---

49

## How to find best $\lambda$ ?

Minimize  $L_{\text{Ridge}}$  in cross validation

→ Hyperparameter tuning

# Penalized Regression

---

50

## LASSO: Least absolute shrinkage and selection operator

Penalize large parameter estimates (L1 regularization)

$$\begin{aligned} L_{\text{LASSO}} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \lambda \sum_{j=1}^m |\beta_j| \\ &= \|y - X\beta\|_2^2 + \lambda \|\beta\|_1 \end{aligned}$$

# Penalized Regression

---

50

## LASSO: Least absolute shrinkage and selection operator

Penalize large parameter estimates (L1 regularization)

$$\begin{aligned} L_{\text{LASSO}} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \lambda \sum_{j=1}^m |\beta_j| \\ &= \|y - X\beta\|_2^2 + \lambda \|\beta\|_1 \end{aligned}$$

No closed-form solution

# Penalized Regression

---

50

## LASSO: Least absolute shrinkage and selection operator

Penalize large parameter estimates (L1 regularization)

$$\begin{aligned} L_{\text{LASSO}} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \lambda \sum_{j=1}^m |\beta_j| \\ &= \|y - X\beta\|_2^2 + \lambda \|\beta\|_1 \end{aligned}$$

No closed-form solution

**Shrink parameter estimates to (exactly) zero**

# Penalized Regression

---

51

## Elastic net: Combination of Ridge and LASSO

L1 and L2 regularization

$$\begin{aligned} L_{\text{Elnet}} &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \lambda_1 \sum_{j=1}^m |\beta_j| + \lambda_2 \sum_{j=1}^m \beta_j^2 \\ &= \|y - X\beta\|_2^2 + \lambda_1 \|\beta\|_1 + \lambda_2 \|\beta\|_2^2 \end{aligned}$$

# Penalized Regression

---

52

## Advantages of penalized regression

- Reduces overfitting
- Avoid multicollinearity issues of (non-penalized) regression models
  - Work well with high-dimensional data
- Same general concept of (non-penalized) regression models
  - Interpretable model
- Better prediction performance than non-penalized regression (less variance)
- Implicit variable selection (LASSO)

# Penalized Regression

---

53

## Disadvantages of penalized regression

- Biased parameter estimates
- Cannot use statistical inference methods used in non-penalized regression
- Interactions and non-linear effects have to be explicitly specified
- Often worse prediction performance than (other) machine learning algorithms

# Outline

---

54


1. Introduction
2. Supervised Learning
3. Decision Trees & Random Forests
4. Model Evaluation & Resampling
5. Penalized Regression
6. Artificial Neural Networks
7. Hyperparameter Tuning & Benchmarking
8. Discussion



# Artificial Neural Networks

---

55




# Artificial Neural Networks

---

56


For each hidden unit:



# Artificial Neural Networks

## Activation function: Sigmoid


57



# Artificial Neural Networks

## Special cases

58



Identity: Linear model

Step: Perceptron

# Artificial Neural Networks

## How to fit a neural network?

59


**Loss function:** Error as function of network weights, e.g,

$$L(\mathcal{W}) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Aim: Find weights that minimize error

## Gradient descent

Adjust weights in direction with steepest descent




# Artificial Neural Networks

## How to fit a neural network?

60

### Backpropagation

1. Initialize weights randomly
2. For  $k$  iterations repeat
  - a) Compute error function
  - b) Adjust weights in output layer
  - c) Propagate error backwards through network and adjust weights




# Artificial Neural Networks

---

Do neural networks overfit?



61



# Artificial Neural Networks

Do neural networks overfit?

61







# Artificial Neural Networks

Do neural networks overfit?



61



# Artificial Neural Networks

Do neural networks overfit?

61



# Artificial Neural Networks

---

62

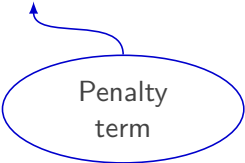
## L2 regularization

$$L(\mathcal{W}) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \lambda \sum_{\mathbf{w} \in \mathcal{W}} \mathbf{w}^2$$

### Alternative: Dropout

Temporarily remove units while fitting

### Other alternative: Early stopping




Penalty  
term

# Artificial Neural Networks

---

## What is (not) deep learning?

63




# Artificial Neural Networks

64

## What is deep learning?

Idea: Build hierarchy of concepts (representations)


## Convolutional neural networks



# Artificial Neural Networks

65

## Example



# Artificial Neural Networks

---

66

## Advantages of neural networks

- Can fit any complicated function almost perfectly
- Learns representations
- Online learning possible
- Prediction very fast (matrix multiplication)
- Can be parallelized (GPU computing)

# Artificial Neural Networks

---

67

## Disadvantages of neural networks

- Prone to overfitting
- Difficult to design good networks
- Interpretation very difficult (black box)
- Learning can be slow
- Statistical properties not well studied



# Outline

---

68

1. Introduction
2. Supervised Learning
3. Decision Trees & Random Forests
4. Model Evaluation & Resampling
5. Penalized Regression
6. Artificial Neural Networks
7. Hyperparameter Tuning & Benchmarking
8. Discussion

# Hyperparameter Tuning

---

69

## Hyperparameters

Learners have hyperparameters, e.g.:

- Number of nearest neighbors  $k$
- Depth of a tree
- Number of features to consider in each split of a random forest (mtry)
- Architecture of neural network

## Most learners have several hyperparameters

Have to be jointly optimized

# Hyperparameter Tuning

---

70

## Search entire parameter space

- All possible combinations
- Grid search
- Randomly select combinations
- Model-based optimization


## Use resampling

- Evaluate each parameter combination on all resampling iterations/folds
- Choose parameter maximizing aggregated performance measure

# Hyperparameter Tuning

## Grid search

71



# Hyperparameter Tuning

---

72

## Grid search

### Advantages

- Easy to implement
- All parameter types possible
- Easily parallelized


### Disadvantages

- Computationally intensive
- Inefficient: Searches large irrelevant areas
- Arbitrary: Which values / discretization?

# Hyperparameter Tuning

## Random search

73



# Hyperparameter Tuning

---

74

## Random search

### Advantages

- Same as grid search: Easy to implement, all parameter types possible, trivial parallelization
- Easy to adjust to computational budget
- No discretization
- Superior performance compared to grid search

### Disadvantages

- Computationally intensive
- Inefficient: Searches large irrelevant areas

# Hyperparameter Tuning

---

75

## Model-based optimization

### Surrogate model

Learn relationship between hyperparameters and prediction performance

### Algorithm

1. Pick initial configuration (e.g. random)
2. Learn surrogate model
3. Predict new configuration with surrogate model
4. Repeat steps 2 and 3



# Hyperparameter Tuning

---

76

## Model-based optimization

### Advantages

- All parameter types possible
- Efficient: Focus on promising areas
- Superior performance compared to grid and random search

### Disadvantages

- Computationally intensive
- Non-trivial parallelization
- Harder to implement

# Benchmarking

---

77

## How can performance be compared?


### Be fair!

- Compare all learners and models on same data
- Tune parameters of all learners
- Don't overfit
- Don't publish over-optimistic results


**Never learn, tune or evaluate on same data!**

# Nested Resampling

78



# Model Selection



# Benchmarking

---

80

## How to build a final model?

1. Select best learner with nested resampling
2. Find optimal hyperparameters of best learner with resampling
3. Train best learner with optimal hyperparameters on full data

# Outline

---

81

1. Introduction
2. Supervised Learning
3. Decision Trees & Random Forests
4. Model Evaluation & Resampling
5. Penalized Regression
6. Artificial Neural Networks
7. Hyperparameter Tuning & Benchmarking
8. Discussion

# Discussion

---

82

Is there a single best learner?

**No!**

## General recommendations

- Typically  $\text{RF} \approx \text{Boosting} > \text{Tree} > \text{kNN}$
- RF robust, easy to tune and fast
- Boosting often slightly better than RF on tabular data (when properly tuned)
- SVM good alternative for binary classification with numerical features (when properly tuned)
- Image, text and speech data  $\rightarrow$  Deep Learning

# Discussion

---

## Important aspects when applying machine learning 83

- Never use default parameter settings
- Tune parameters!
- Tune parameters jointly!
- Parameter tuning simple and straightforward for
  - kNN
  - Decision trees
  - Boosting
  - Random forests
- Parameter tuning complex and not straightforward for
  - SVM: parameters depend on kernel
  - ANN: tuning of architecture
- Use adequate resampling strategy
- Gold standard: nested cross validation